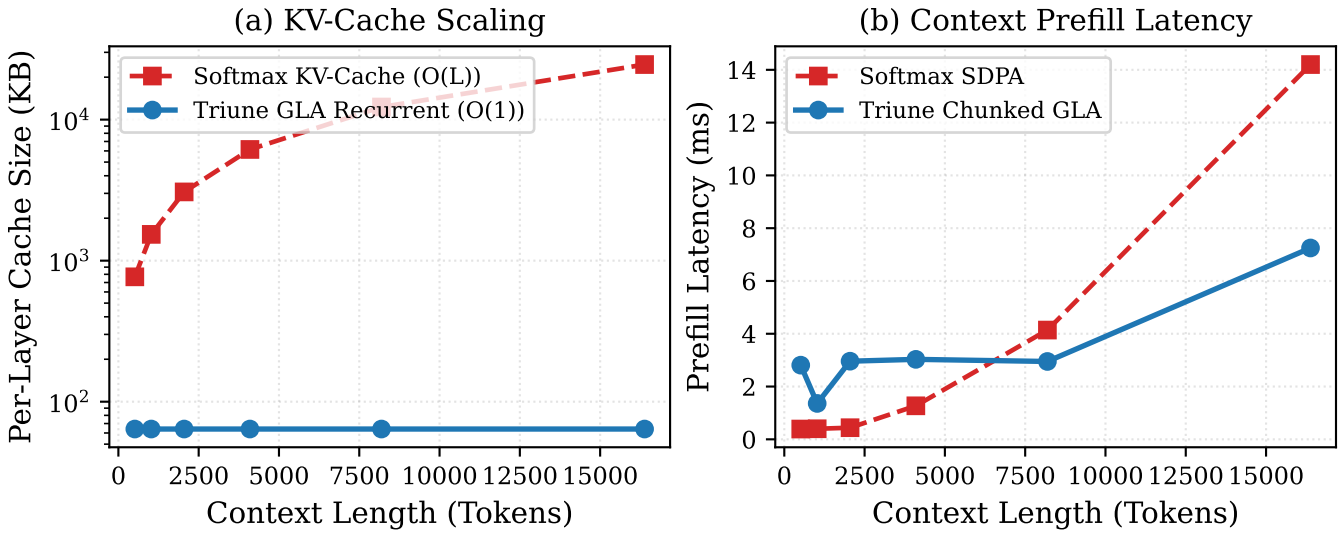
(a) KV-Cache Scaling
Softmax KV-Cache (O(L))
Triune GLA Recurrent (O(1))
Per-Layer Cache Size (KB)
10^4
10^3
10^2
0
2500
5000
7500
10000
12500
15000
Context Length (Tokens)
(b) Context Prefill Latency
Softmax SDPA
Triune Chunked GLA
Prefill Latency (ms)
0
2
4
6
8
10
12
14
0
2500
5000
7500
10000
12500
15000
Context Length (Tokens)